



Praveen Kumar Chittem

AI ENABLEMENT &
FORWARD DEPLOYMENT LEAD

CONTACT

✉ impravin22@gmail.com

☎ +886 958 605 515

📍 Taipei, Taiwan

ENABLEMENT

AI Adoption

workflow discovery · reusable AI skills and prompt templates · AI-fluency framework · adoption metrics

Claude & Anthropic

Claude corporate workspaces · Claude Code · Anthropic API · MCP server design · AI code review

Google Gemini

Gemini API · Gemini Vision · structured output

Governance

AI use policy · consent and PII scrubbing · OWASP LLM Top 10 · SSO and role-based access

STACK

Python · TypeScript · SQL · FastAPI · Next.js · REST APIs · LangGraph · DSPy · PostgreSQL / pgvector · Dagster · Celery / Redis · AWS · Docker

SELECTED IMPACT

1,300+ platform users, sole builder

9× faster goal updates

25 exec sessions run over MCP

3 eng pods on AI code review

LANGUAGES

English (professional) · Mandarin (learning)

PROFILE

AI engineer who sits between the business and the build: I work with department heads to find where work gets stuck, turn it into Claude and Gemini solutions people use daily, and wrap them in guardrails that security and legal sign off on. Six years shipping production AI end to end; now leading the forward-deployed AI pod in ViewSonic's CEO Office.

EXPERIENCE

Forward Deployment Pod Lead

Jul 2026 – Present

ViewSonic · CEO Office | Taipei, Taiwan

- **Discovery to delivery.** Repeatable intake: find bottlenecks in org-wide workflow data, prioritise, validate the process with the affected team, get a go/no-go, then build. Spans marketing, supply chain, legal and the EDU business unit.
- **Claude and MCP rollout.** Built the platform's MCP server so staff query goals, AI-maturity and org data in Claude; turned a meeting framework into an MCP workflow that runs a 25-session leadership meeting in London this October.
- **Agentic delivery.** Built an automated issue solver for an internal pricing tool I own: an agent picks up the issue, writes and tests the fix, clears every pipeline gate and releases to dev on its own; production still needs a human approval.
- **Developer enablement.** AI-assisted code review across three engineering pods (local, pre-PR and AI-bot GitHub workflows); wrote the business case for Claude seats over uncapped cloud model spend.
- **Governance.** Drafted the data-privacy rule for a workflow-analytics pilot on the company's Data & AI Use Policy: individual consent and PII scrubbing before any email data reaches a model; SSO and role-based access on every tool.
- **Programme.** Several workstreams in parallel; client lead and technical coach for a five-student Northwestern University graduate capstone (Sep to Dec 2026).

Senior AI Systems Architect

Mar 2022 – Jun 2026

ViewSonic · CEO Office | Taipei, Taiwan

- **Built the internal AI platform** from a blank repo to 1,300+ users: goal-setting and OKR coaching, performance evaluation and workforce AI-maturity.
- **Adoption.** Co-developed the AI-fluency framework and an AI-maturity assistant that rates every employee L1 to L5; org-wide adoption dashboards.
- **Reusable AI patterns.** One template (conversation, structured scoring, deterministic aggregation, narrative) reused across modules on DSPy signatures. Rule: *prompts for judgment, code for determinism*.
- **Safe by default.** Full OWASP LLM Top 10 coverage, per-user answer redaction at inference, SQL validation on Text-to-SQL, versioned ethics gate.
- **Gemini in production.** Gemini Vision pipeline turning PPTX, DOCX, PDF and images into brand-compliant decks; an auditor loop remakes weak output.
- **Earlier ML.** Distilled a 600M-parameter model to 2× faster at parity; ran an 8-GPU Kubernetes cluster for distributed PyTorch training.

AI Software Developer

Aug 2020 – Mar 2022

Neurelli | Taipei, Taiwan

- Shipped industrial computer vision to production (licence-plate recognition on TensorRT/CUDA, LCD defect detection), from data to deployment.

EDUCATION

M.Sc. Electrical Engineering & Computer Science

2018 – 2020

National Taipei University of Technology | Taiwan

- Thesis: CUDA-accelerated flood simulation, 15× speedup; published at IEEE.